

NON CLUSTERED INDEX GUIDE

- A nonclustered index contains the index **key values and row locators** that **point to the storage location** of the table data
- As discussed, we can have **multiple nonclustered indexes** on a table or indexed view
- Nonclustered indexes should be designed for **frequently used queries**
- Nonclustered indexes are optimal choice for **EXACT match** queries because the index contains entries describing the exact location in the table
- Those tables with **low update** requirements -databases that contain heavily updated tables should avoid over-indexing.
- **Read-only data** can benefit from many nonclustered indexes
- OLTP applications and large numbers of indexes on a table affect the performance of INSERT, UPDATE, DELETE, and MERGE statements **BECAUSE all indexes must be adjusted** appropriately as data in the table changes
- Use JOIN or GROUP BY clauses
- Create multiple nonclustered indexes on columns involved in join and grouping operations
- Create a non cluster **a column(s) frequently** involved in search conditions of a query using **WHERE clause**
- Create a non cluster **a column(s)** to **cover the query** (meaning if the where clause has more than one predicate, cover those columns)
- When the query optimizer can locate all the column values within the index; table or clustered index data is not accessed resulting in fewer disk I/O operations.
- Create non clustered on **distinct values**, such as a combination of last name and first name
- other